

NAIC AI Governance

Readiness Checklist — 2026

10 questions every insurance compliance team should be able to answer for the NAIC AI Systems Evaluation Tool — now live in 12 states.

Why this matters now: The NAIC AI Systems Evaluation Tool pilot launched March 2026 across 12 states. Examiners are sending structured questionnaires to carriers today. Full nationwide adoption expected at the NAIC Fall Meeting — November 2026. This checklist maps to the four exhibits examiners will request.

How to score yourself: Answer each question honestly. Count your YES answers at the end to see your readiness level.

EXHIBIT A — AI System Inventory

Examiners will ask: What AI systems are you using and what decisions do they influence?

Q1

Can you produce a complete inventory of every AI system in use across your organization — including vendor-embedded models — within 24 hours of an examiner request?

■ Yes — we have a maintained AI system registry with all systems documented ■ Partial — we have a list but it may be incomplete or outdated ■ No — we would need to compile this manually from scratch

Q2

For each AI system, do you have documented records of: the function it performs, which decisions it influences, whether it is vendor-built or in-house, the date it was deployed, and when it was last reviewed?

■ Yes — all fields documented and current for every AI system ■ Partial — some systems documented, others not ■ No — this information is scattered or unavailable

EXHIBIT B — Governance Framework

Examiners will ask: How do you govern AI — oversight structures, accountability, audit trails, policy compliance evidence?

Q3

For any AI decision made in the last 24 months — can you identify which version of your internal compliance policy was active at the exact moment that decision was made?

■ Yes — we have timestamped policy versions linked to every AI decision ■ Partial — we have some records but significant gaps exist ■ No — we would need to reconstruct this from logs and emails

Q4	If your compliance policies changed while an AI workflow was running, do you have evidence that the AI system received and applied the updated rules — not the previous version?
	☐ Yes — we have documented verification of policy updates to AI systems ☐ Partial — we communicate updates but don't formally verify receipt ☐ No — we assume the AI got the update but cannot prove it
Q5	Do you have a continuous audit trail — not just a periodic snapshot — showing AI decision compliance over time, including any instances where stale or outdated policy rules were applied?
	☐ Yes — continuous monitoring with documented drift incidents and resolution ☐ Partial — periodic reviews but not continuous monitoring ☐ No — we only review during examination preparation
Q6	When a human reviewer overrides an AI recommendation, is that override logged — including who made the decision, why they overrode it, and under which policy version — in a tamper-evident record?
	☐ Yes — override logging with reasons is systematic and fully documented ☐ Partial — some overrides logged but not consistently ☐ No — overrides are tracked informally or not at all

EXHIBIT C — High Risk AI System Details

Examiners will ask: For AI systems affecting policyholders — show us model details, validation, and performance documentation.

Q7	For each high-risk AI system (underwriting, claims, pricing, fraud detection) — do you have current documentation of model design, training data sources, validation results, and performance metrics?
	☐ Yes — full model documentation maintained and current for all high-risk systems ☐ Partial — documentation exists but may be incomplete or outdated ☐ No — this documentation does not currently exist
Q8	Have your high-risk AI systems been tested for potential bias or disparate impact — and is that testing documented with results and any remediation actions taken?
	☐ Yes — bias testing completed, documented, and remediation tracked ☐ Partial — some testing done but not systematic or fully documented ☐ No — bias testing has not been formally conducted

EXHIBIT D — AI Data Details

Examiners will ask: What data feeds your AI, where does it come from, and how do you ensure it doesn't create discriminatory outcomes?

Q9

Can you document the data sources feeding each high-risk AI system — including data quality controls, how representative the training data is, and whether any data sources could act as proxies for protected characteristics?

■ Yes — data lineage documented for all high-risk AI systems ■ Partial — some data sources documented but not all ■ No — data lineage documentation does not currently exist

The Examiner's Hardest Question

Q10

If a state examiner called you today and asked: 'Show me every AI-assisted underwriting decision made in the last 6 months, which compliance policy version governed each one, and flag any where the policy had changed but the AI may not have received the update' — how long would it take to produce that answer?

■ Under 1 hour — we have an automated audit trail covering exactly this ■ 1 to 3 days — we can reconstruct it but it requires significant manual effort ■ 1 to 2 weeks — substantial engineering work required to find the answer ■ We cannot answer this question today

Your Score

Score	What it means	What to do next
9-10 Yes	Examination ready. Strong governance posture.	Document everything. Schedule a practice run before your next examination.
6-8 Yes	Partially ready. Gaps exist that examiners will find.	Prioritize the red questions. Build documentation for Exhibit B gaps first.
3-5 Yes	Significant exposure. Examination will be uncomfortable.	Start with AI System Registry and policy version audit trail immediately.
0-2 Yes	High risk. Examination findings likely.	Address before November 2026 nationwide adoption. Start today.

What PolicyProof Covers

PolicyProof produces the Exhibit B documentation examiners request — automatically and continuously. Pre-loaded with insurance demo data. No signup required.

Question	Exhibit	PolicyProof Today
----------	---------	-------------------

Q1 — AI System Inventory	A	✓ AI System Registry — built
Q2 — AI System Documentation	A	✓ All fields captured in registry
Q3 — Policy version per decision	B	✓ Core feature — built
Q4 — Policy update verification	B	✓ Stale drift detection — built
Q5 — Continuous audit trail	B	✓ Governance health timeline — built
Q6 — Override logging	B	Building now — next sprint
Q7 — Model documentation	C	Partial — policy traceability covered
Q8 — Bias testing	C	Out of scope — specialist tool needed
Q9 — Data lineage	D	Out of scope — data governance tool
Q10 — Examiner hard question	All	✓ One-click audit report — built

See PolicyProof in action. Pre-loaded with insurance demo data — underwriting decisions, policy mutations, stale drift detection, and one-click NAIC-formatted audit report. No signup required.

policyproof.vercel.app/app.html

Free resource from PolicyProof — policyproof.vercel.app — May 2026